


# Relational Data with dplyr



# Outline

Intro to relational data

# Outline

Intro to relational data

Mutating Joins

# Outline

Intro to relational data

Mutating Joins

Filtering Joins

# Outline

Intro to relational data

Mutating Joins

Filtering Joins

Set Operations

# Relational Data

Relational data - multiple tables of data that share some common variables / attributes (relations and not individual datasets that are important)

# Relational Data

Relational data - multiple tables of data that share some common variables / attributes (relations and not individual datasets that are important)

A relation will always be between a pair of tables.

# Relational Data

Relational data - multiple tables of data that share some common variables / attributes (relations and not individual datasets that are important)

A relation will always be between a pair of tables.

Subsequent relations will build off of relations of pairs.



# Relational Data

Relational data - multiple tables of data that share some common variables / attributes (relations and not individual datasets that are important)

A relation will always be between a pair of tables.

Subsequent relations will build off of relations of pairs.

Verbs to work with pairs of tables:

# Relational Data

Relational data - multiple tables of data that share some common variables / attributes (relations and not individual datasets that are important)

A relation will always be between a pair of tables.

Subsequent relations will build off of relations of pairs.

Verbs to work with pairs of tables:

Mutating joins - add new variables from matching observations in another

# Relational Data

Relational data - multiple tables of data that share some common variables / attributes (relations and not individual datasets that are important)

A relation will always be between a pair of tables.

Subsequent relations will build off of relations of pairs.

Verbs to work with pairs of tables:

Mutating joins - add new variables from matching observations in another

Filtering joins - filter observations from one dataset based on whether or not they match an observation in the other table

# Relational Data

Relational data - multiple tables of data that share some common variables / attributes (relations and not individual datasets that are important)

A relation will always be between a pair of tables.

Subsequent relations will build off of relations of pairs.

Verbs to work with pairs of tables:

Mutating joins - add new variables from matching observations in another

Filtering joins - filter observations from one dataset based on whether or not they match an observation in the other table

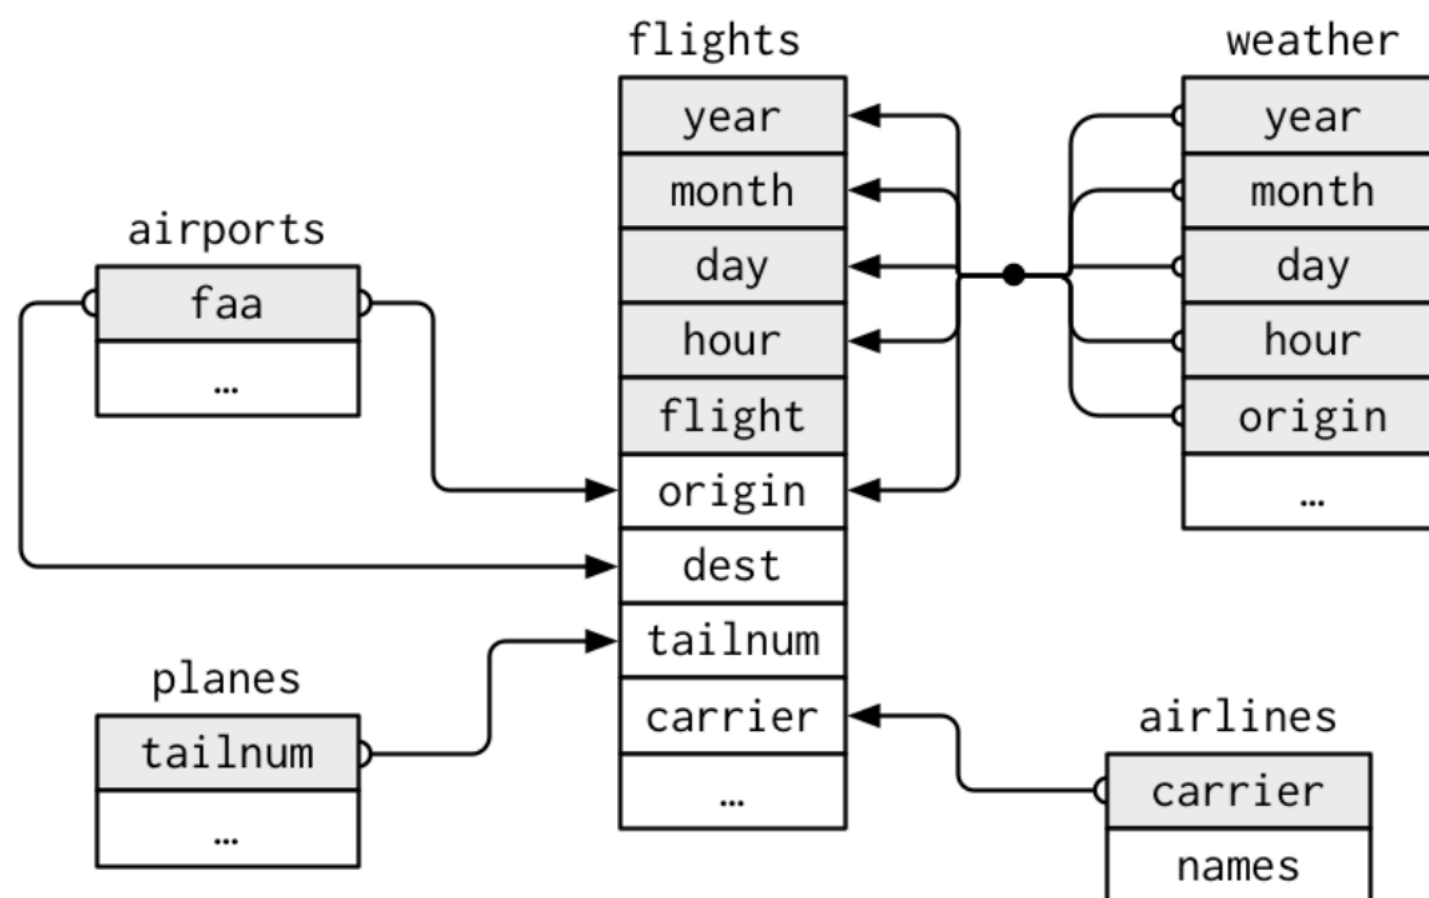
Set operations - treats observations like set elements

# nycflights13

The *nycflights13* R package contains similar info as the Houston flights data but also contains data frames about the airlines, weather, and airports.

# nycflights13

The *nycflights13* R package contains similar info as the Houston flights data but also contains data frames about the airlines, weather, and airports.



# Keys

The variables we use to connect the datasets are called keys

# Keys

The variables we use to connect the datasets are called keys

Three types of keys:



# Keys

The variables we use to connect the datasets are called keys

Three types of keys:

- Primary Key - uniquely identifies an observation in its own table

# Keys

The variables we use to connect the datasets are called keys

Three types of keys:

- Primary Key - uniquely identifies an observation in its own table

- Foreign Key - uniquely identifies an observation in another table

# Keys

The variables we use to connect the datasets are called keys

Three types of keys:

- Primary Key - uniquely identifies an observation in its own table

- Foreign Key - uniquely identifies an observation in another table

- A variable can be both a primary and foreign key!

# Keys

The variables we use to connect the datasets are called keys

Three types of keys:

- Primary Key - uniquely identifies an observation in its own table

- Foreign Key - uniquely identifies an observation in another table

- A variable can be both a primary and foreign key!

- Surrogate Key - a key that is created because the dataset lacks a primary key

# Keys

The variables we use to connect the datasets are called keys

Three types of keys:

- Primary Key - uniquely identifies an observation in its own table

- Foreign Key - uniquely identifies an observation in another table

- A variable can be both a primary and foreign key!

- Surrogate Key - a key that is created because the dataset lacks a primary key

A primary key and the corresponding foreign key in another table for the *relation*.

# Mutating Joins

Setup:

# Mutating Joins

Setup:

```
x <- tribble(
  ~key, ~val_x,
  1, "x1",
  2, "x2",
  3, "x3"
)
```

```
y <- tribble(
  ~key, ~val_y,
  1, "y1",
  2, "y2",
  4, "y3"
)
```

# Mutating Joins

x		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y3

Setup:

```
x <- tribble(
  ~key, ~val_x,
  1, "x1",
  2, "x2",
  3, "x3"
)

y <- tribble(
  ~key, ~val_y,
  1, "y1",
  2, "y2",
  4, "y3"
)
```



# Mutating Joins


x		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y3

Setup:

```
x <- tribble(
  ~key, ~val_x,
  1, "x1",
  2, "x2",
  3, "x3"
)
```

```
y <- tribble(
  ~key, ~val_y,
  1, "y1",
  2, "y2",
  4, "y3"
)
```

A join is a way of connecting the rows in x to 0,1, or more rows in y



# Mutating Joins


x		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y3

`inner_join()` - matches pairs of observations whenever their keys are equal (drops everything else)

# Mutating Joins

	x		y
1	x1	1	y1
2	x2	2	y2
3	x3	4	y3


`inner_join()` - matches pairs of observations whenever their keys are equal (drops everything else)



# Mutating Joins

x		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y3

`inner_join()` - matches pairs of observations whenever their keys are equal (drops everything else)



```
inner_join(x, y, by = "key")
```

```
x %>%
```

```
inner_join(y, by = "key")
```

# Mutating Joins

Outer joins (`left_join()`, `right_join()`, `full_join()`) - keeps observations that appear in at least one of the tables

# Mutating Joins

Outer joins (`left_join()`, `right_join()`, `full_join()`) - keeps observations that appear in at least one of the tables

`left_join()` - keeps all observations in x (left dataset)

# Mutating Joins

Outer joins (`left_join()`, `right_join()`, `full_join()`) - keeps observations that appear in at least one of the tables

`left_join()` - keeps all observations in x (left dataset)

`right_join()` - keeps all observations in y (right dataset)

# Mutating Joins

Outer joins (`left_join()`, `right_join()`, `full_join()`) - keeps observations that appear in at least one of the tables

`left_join()` - keeps all observations in x (left dataset)

`right_join()` - keeps all observations in y (right dataset)

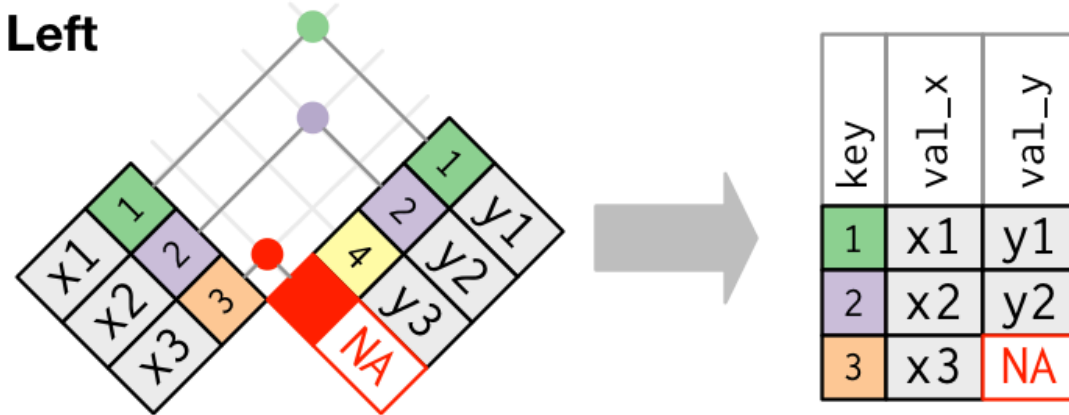
`full_join()` - keeps all observations in both datasets



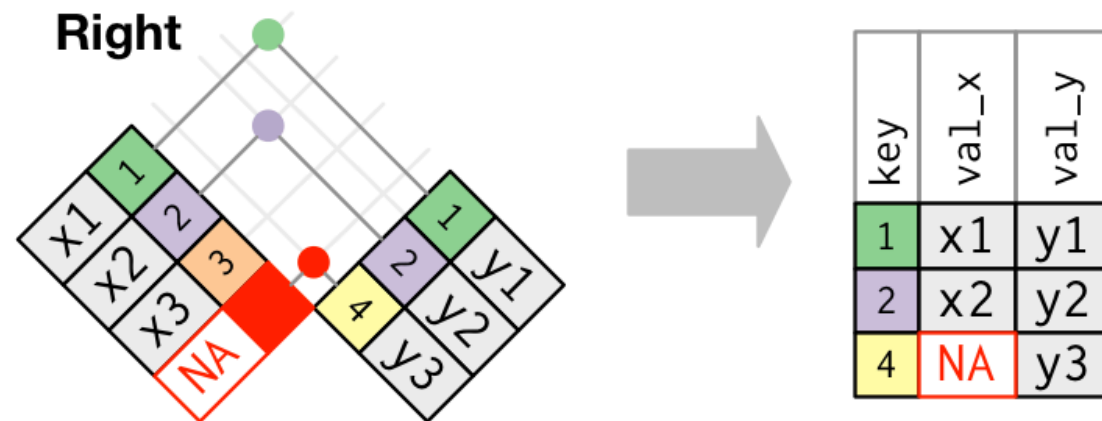
# Mutating Joins

	x	y
1	x1	y1
2	x2	y2
3	x3	
		4
		y3

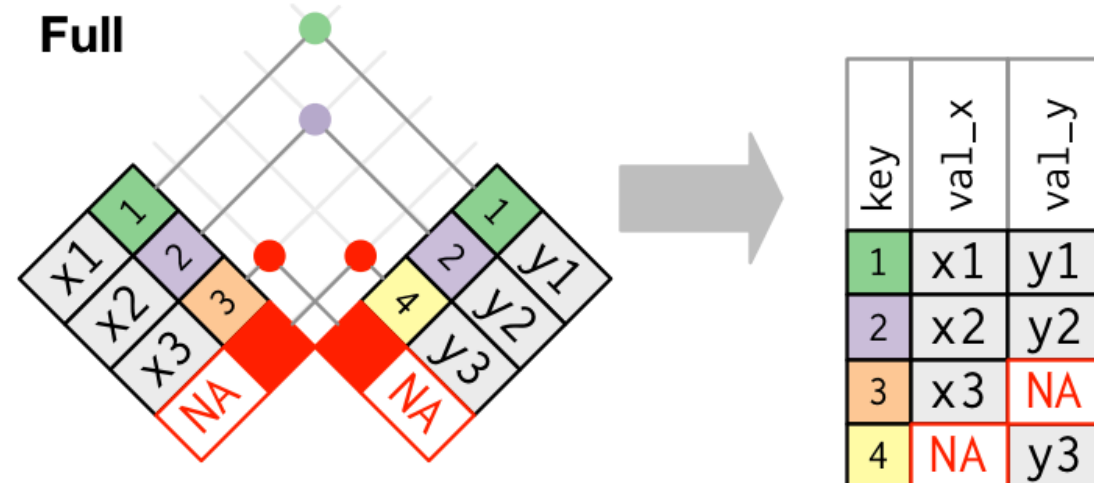
Left



Right



Full




# Mutating Joins

Duplicate keys - When the keys are not unique.


# Mutating Joins

Duplicate keys - When the keys are not unique.



# Mutating Joins


Duplicate keys - When the keys are not unique.




When you join duplicate keys, you get all possible combos

# Mutating Joins

Duplicate keys - When the keys are not unique.



When you join duplicate keys, you get all possible combos



# Defining Key Columns

There are multiple ways to connect tables with the “by” arg.

# Defining Key Columns

There are multiple ways to connect tables with the “by” arg.

by = **NULL** - uses all variables in both datasets (natural join)

# Defining Key Columns

There are multiple ways to connect tables with the “by” arg.

by = **NULL** - uses all variables in both datasets (natural join)

by = **"x"** - a character vector. Uses only variables defined in key



# Defining Key Columns

There are multiple ways to connect tables with the “by” arg.

`by = NULL` - uses all variables in both datasets (natural join)

`by = "x"` - a character vector. Uses only variables defined in key

`by = c("a" = "b")` - named character vector. Will match variable *a* in table *x* to variable *b* in table *y*

# Filtering Joins

`semi_join()` and `anti_join()`

# Filtering Joins

`semi_join()` and `anti_join()`

`semi_join()` - keeps all observations in x that have a match in y

# Filtering Joins

`semi_join()` and `anti_join()`

`semi_join()` - keeps all observations in x that have a match in y

`anti_join()` - drops all observations in x that have a match in y

# Filtering Joins

`semi_join()` and `anti_join()`

`semi_join()` - keeps all observations in x that have a match in y

`anti_join()` - drops all observations in x that have a match in y

`semi_join()`:



# Filtering Joins

`semi_join()` and `anti_join()`

`semi_join()` - keeps all observations in x that have a match in y

`anti_join()` - drops all observations in x that have a match in y

`semi_join()`:





# Filtering Joins

`semi_join()` and `anti_join()`

`semi_join()` - keeps all observations in x that have a match in y

`anti_join()` - drops all observations in x that have a match in y

`semi_join()`:



`anti_join()`:


# Filtering Joins

`semi_join()` and `anti_join()`


`semi_join()` - keeps all observations in x that have a match in y

`anti_join()` - drops all observations in x that have a match in y

`semi_join()`:



`anti_join()`:





# Set Operations

Set operations can be helpful when you break down a simple complex filter into simpler filters

# Set Operations

Set operations can be helpful when you break down a simple complex filter into simpler filters

`intersect(x, y)` - returns only observations in both x and y

# Set Operations

Set operations can be helpful when you break down a simple complex filter into simpler filters

`intersect(x, y)` - returns only observations in both x and y

`union(x, y)` - returns unique observations in x and y

# Set Operations

Set operations can be helpful when you break down a simple complex filter into simpler filters

`intersect(x, y)` - returns only observations in both x and y

`union(x, y)` - returns unique observations in x and y

`setdiff(x, y)` - returns observations in x, but not y