

# Relational Data with dplyr



# Outline

Intro to relational data

Mutating Joins

Filtering Joins

Set Operations

# Relational Data

Relational data - multiple tables of data that share some common variables / attributes (relations and not individual datasets that are important)

A relation will always be between a pair of tables.

Subsequent relations will build off of relations of pairs.

Verbs to work with pairs of tables:

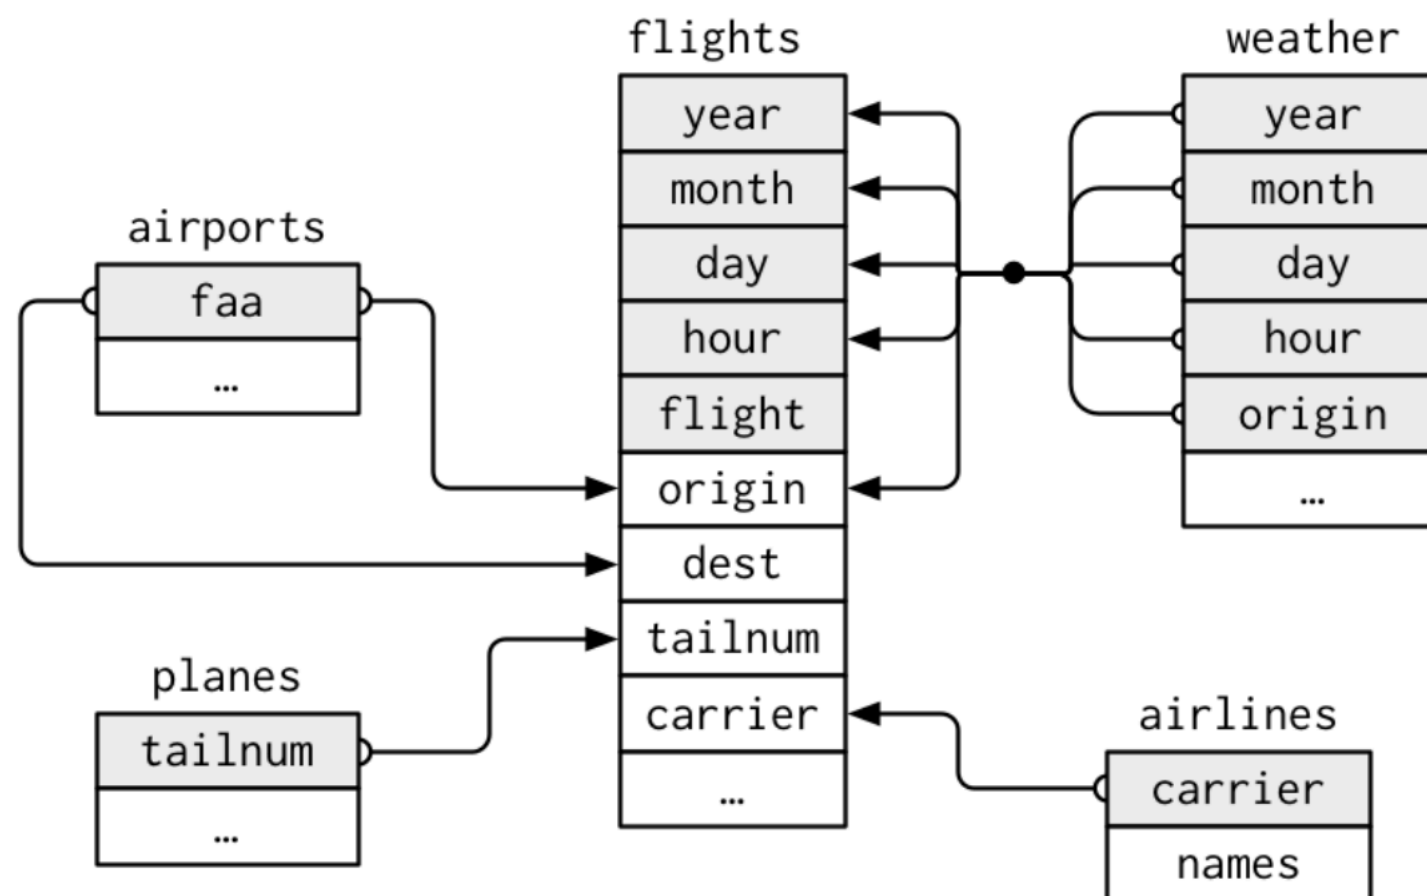
Mutating joins - add new variables from matching observations in another

Filtering joins - filter observations from one dataset based on whether or not they match an observation in the other table

Set operations - treats observations like set elements

# nycflights13

The *nycflights13* R package contains similar info as the Houston flights data but also contains data frames about the airlines, weather, and airports.



# Keys

The variables we use to connect the datasets are called keys

Three types of keys:

- Primary Key - uniquely identifies an observation in its own table

- Foreign Key - uniquely identifies an observation in another table

- A variable can be both a primary and foreign key!

- Surrogate Key - a key that is created because the dataset lacks a primary key

A primary key and the corresponding foreign key in another table form the *relation*.

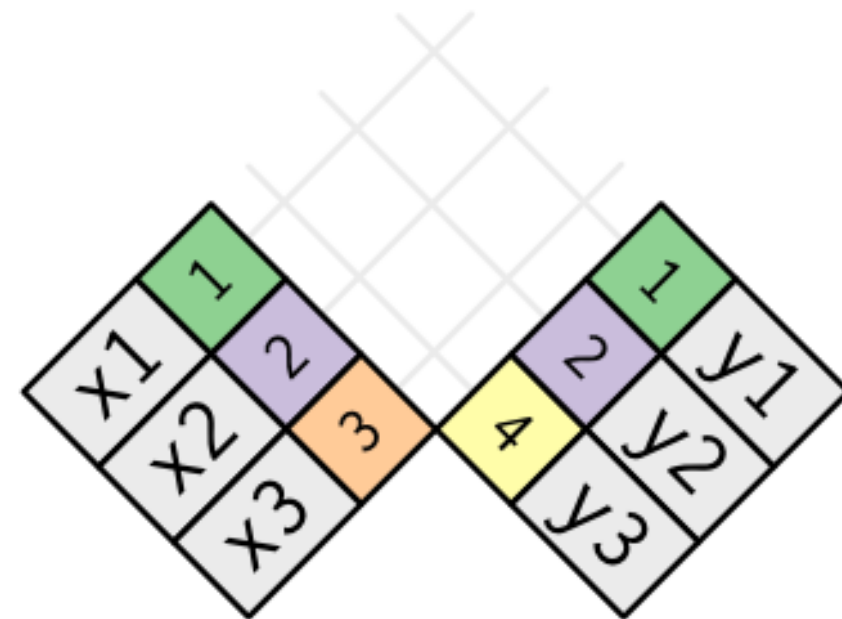
# Mutating Joins

| x |    | y |    |
|---|----|---|----|
| 1 | x1 | 1 | y1 |
| 2 | x2 | 2 | y2 |
| 3 | x3 | 4 | y3 |

Setup: `x <- tribble(  
 ~key, ~val_x,  
 1, "x1",  
 2, "x2",  
 3, "x3"  
)`

`y <- tribble(  
 ~key, ~val_y,  
 1, "y1",  
 2, "y2",  
 4, "y3"  
)`


A join is a way of connecting the rows in x to 0,1, or more rows in y



# Mutating Joins

| x |    | y |    |
|---|----|---|----|
| 1 | x1 | 1 | y1 |
| 2 | x2 | 2 | y2 |
| 3 | x3 | 4 | y3 |

`inner_join()` - matches pairs of observations whenever their keys are equal (drops everything else)



```
inner_join(x, y, by = "key")
```

```
x %>%
```

```
inner_join(y, by = "key")
```

# Mutating Joins

Outer joins (`left_join()`, `right_join()`, `full_join()`) - keeps observations that appear in at least one of the tables

`left_join()` - keeps all observations in x (left dataset)

`right_join()` - keeps all observations in y (right dataset)

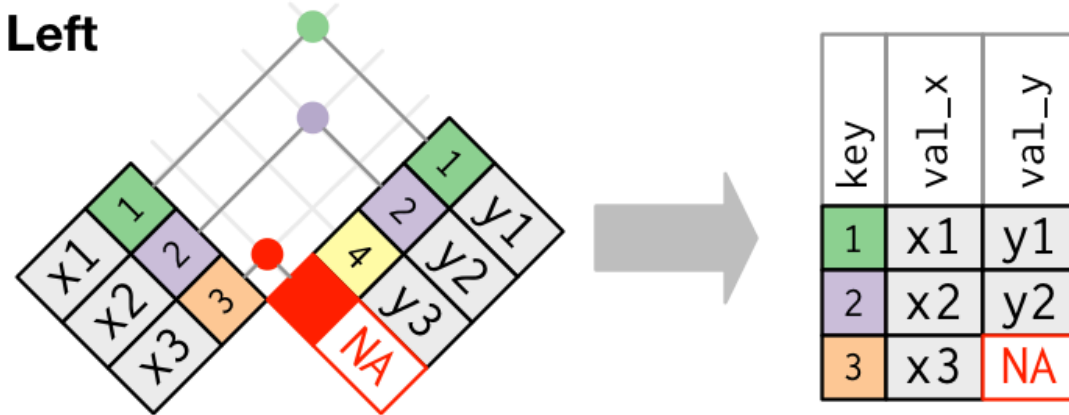
`full_join()` - keeps all observations in both datasets



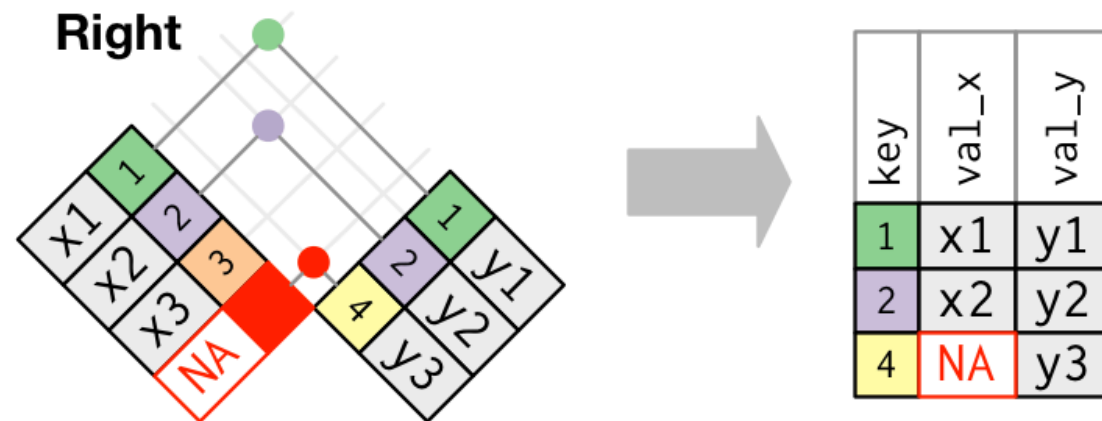
# Mutating Joins

|   | x  | y  |
|---|----|----|
| 1 | x1 | y1 |
| 2 | x2 | y2 |
| 3 | x3 |    |
|   |    | 4  |
|   |    | y3 |

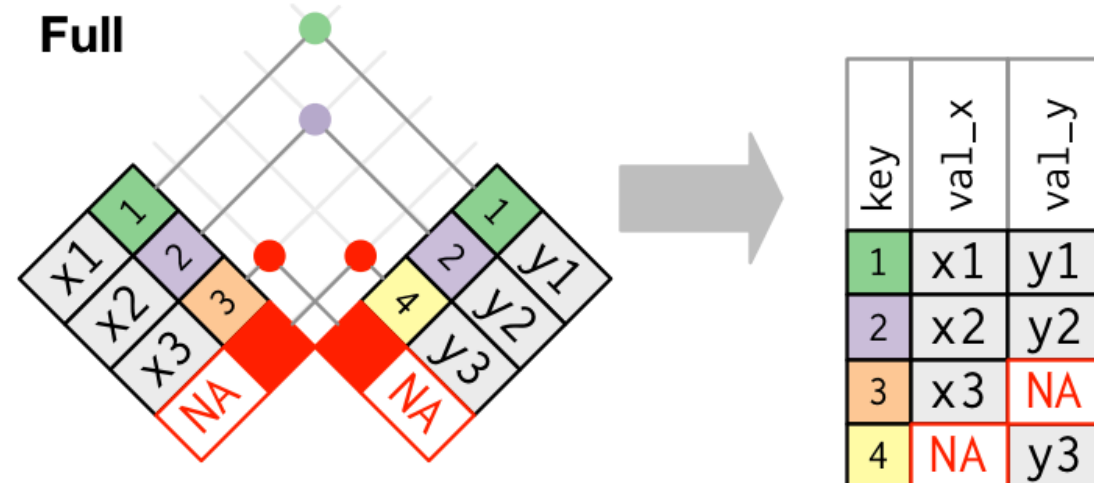
Left



Right

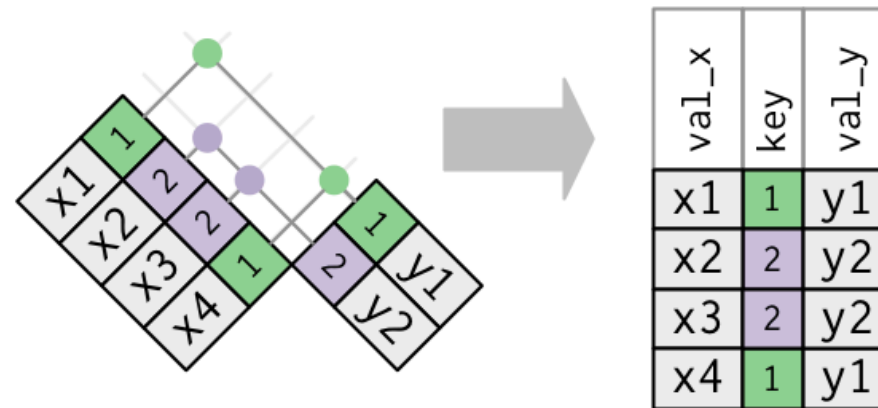


Full

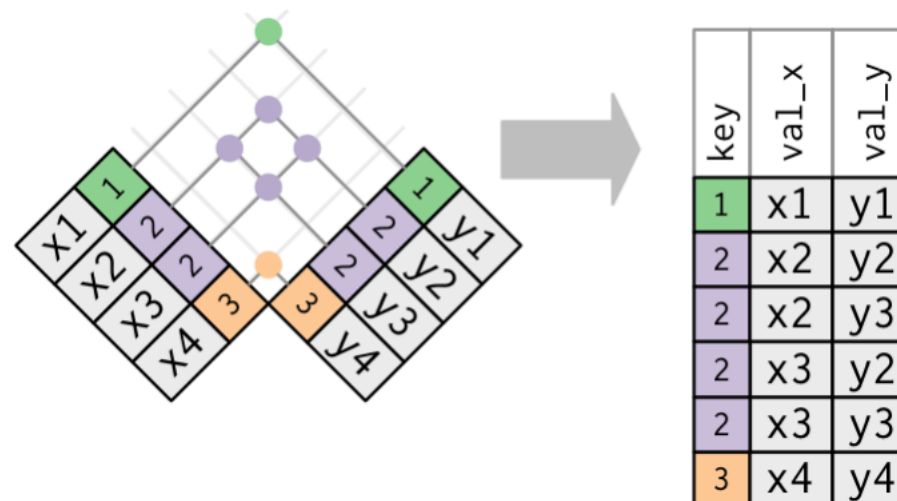


# Mutating Joins

Duplicate keys - When the keys are not unique.



When you join duplicate keys, you get all possible combos



# Defining Key Columns

There are multiple ways to connect tables with the “by” arg.

`by = NULL` - uses all variables in both datasets (natural join)

`by = "x"` - a character vector. Uses only variables defined in key

`by = c("a" = "b")` - named character vector. Will match variable *a* in table *x* to variable *b* in table *y*

`by = join_by(a == b)` - The `join_by()` function is a new function that specifies all types of joins. This will match variable *a* in table *x* to variable *b* in table *y* as well

# Your Turn!

The dplyr package has two datasets called *band\_members* and *band\_instruments* that contain information about people from famous bands.

Use a natural join with the 4 different mutating join functions to see how the output differs.


# Filtering Joins

semi\_join() and anti\_join()


semi\_join() - keeps all observations in x that have a match in y

anti\_join() - drops all observations in x that have a match in y

semi\_join():



anti\_join():



# Your Turn!

The dplyr package has two datasets called *band\_members* and *band\_instruments* that contain information about people from famous bands.

Use a natural join with the 2 different filtering join functions to see how the output differs.

# Set Operations

Set operations can be helpful when you break down a single complex filter into simpler filters

`intersect(x, y)` - returns only observations in both x and y

`union(x, y)` - returns unique observations in x and y

`setdiff(x, y)` - returns observations in x, but not y